

# **Convolutional Neural Networks & Weights and Biases (WandB)**

**Antonio Fonseca**

# Agenda



## 1) Quick recap

- Regularization
- Capacity, Overfitting and Underfitting
- Debugging tips
- Family of optimizers

## 2) Convolutional Neural Networks

- Spatial locality structure
- Kernels, padding, pooling
- Classification tasks
- Saliency Analysis
- Tutorial: data batching, classification of satellite images, WandB

# Regularization




- Prefers to share smaller weights
- Makes model smoother
- More Convex


# Extra Regularization for Neural Nets

Dropout: accuracy in the absence of certain information


- Prevent dependence on any one (or any small combination) of neurons




(a) Standard Neural Net



(b) After applying dropout.




(a) At training time




(b) At test time



# Expectation

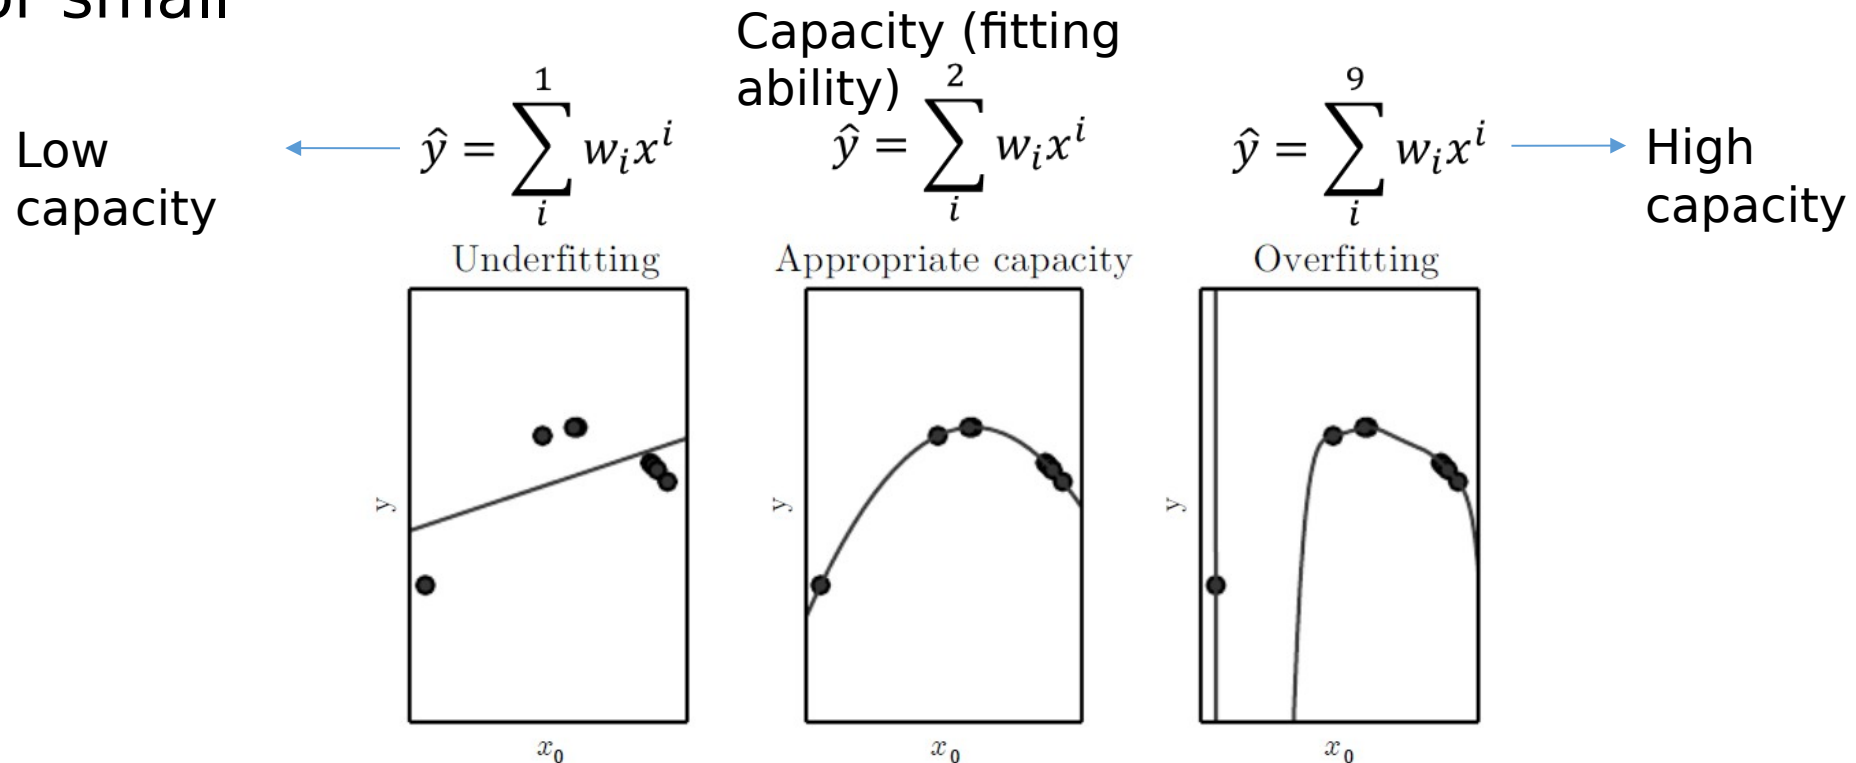


# Reality



# Capacity, Overfitting and Underfitting

- 1) Make training error small
- 2) ~~Make~~ Make the gap between training and test error small



# How training works

1. In each ***epoch***, randomly shuffle the training data
2. Partition the shuffled training data into ***mini-batches***
3. For each mini-batch, apply a single step of **gradient descent**
  - **Gradients** are calculated via ***backpropagation***
4. Train for multiple epochs


# Debugging a neural network

- What can we do?
  - Should we change the learning rate?
  - Should we initialize differently?
  - Do we need more training data?
  - Should we change the architecture?
  - Should we run for more epochs?
  - Are the features relevant for the problem?
- Debugging is an art
  - We'll develop good heuristics for choosing good architectures and hyper parameters (or use tools to help with that)



# Perceptron: Threshold Logic


$$\mathcal{L}_{\text{perc}}(\mathbf{x}, y) = \begin{cases} 0 & \text{if } y\mathbf{w}^\top \mathbf{f}(\mathbf{x}) > 0 \\ -y\mathbf{w}^\top \mathbf{f}(\mathbf{x}) & \text{if } y\mathbf{w}^\top \mathbf{f}(\mathbf{x}) \leq 0 \end{cases}$$



# Activation functions


## Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




## tanh

$$\tanh(x)$$




## ReLU

$$\max(0, x)$$



## Leaky ReLU

$$\max(0.1x, x)$$




## Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$


## ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$




# Limitations of the Perceptron








AND		
$I_1$	$I_2$	out
0	0	0
0	1	0
1	0	0
1	1	1




OR		
$I_1$	$I_2$	out
0	0	0
0	1	1
1	0	1
1	1	1



Perceptr  
on



# Architecture of Neural Networks



But how do we  
train it?

- Sometimes called multi-layer perceptron (MLP)
- Output from one layer is used as input for the next (Feedforward network)

# Optimizers


## Hyperparameters

- Learning rate ( $\alpha$ )


$$\Delta w_k = -\alpha \frac{\partial E}{\partial w_k}$$
$$= -\alpha \frac{\partial}{\partial w_k} \left( \frac{1}{m} \sum_i (w^T X_i - y_i)_i^2 \right)$$

$$w_{i+1} = w_i + \Delta w_k$$

## Stochastic gradient descent (**SGD**)



Local  
Minima



Multiple  
samples




# Optimizers

Hyperparameters

- Learning rate ( $\alpha$ )
- Momentum ( $\beta$ )

$$v_{i+1} = v\beta - \alpha \frac{\partial}{\partial w_k} \left( \frac{1}{m} \sum_i (w^T X_i - y_i)^2 \right)$$

$$w_{i+1} = w_i + v$$



SG

SGD+Momentum

Stochastic gradient descent with momentum  
(**SGD+Momentum**)

# Optimizers

Hard to pick right hyperparameters

- Small learning rate: long convergence time
- Large learning rate: convergence problems

**Adagrad**: adapts learning rate to each parameter

$$\Delta w_{k,t} = -\alpha \frac{\partial E}{\partial w_{k,t}} = -\alpha \nabla_w E(w_t)$$

- Learning rate might decrease too fast
- Might not converge

$$g_{t,i} = \nabla_w E(w_{t,i})$$

$$G_{t+1,i} = G_{t,i} + g_{t,i} \odot g_{t,i}$$



$$w_{t+1,i} = w_{t,i} - \frac{\alpha}{\sqrt{G_{t,i} + \epsilon}} g_{t,i}$$

# Optimizers

RMSprop: decaying average of the past squared gradients

$$E[g^2]_t = \gamma E[g^2]_{t-1} + (1 - \gamma) g_t^2$$

→ Decaying average

$$\Delta w_{k,t} = -\alpha \frac{\partial E_t}{\partial w_{k,t}} = -\alpha \nabla_w E(w_t) = -\alpha g_{t,i}$$

$$g_{t,i} = \nabla_w E(w_{t,i})$$

$$G_{t+1,i} = \gamma G_{t,i} + (1 - \gamma) g_{t,i} \odot g_{t,i}$$



$$w_{t+1,i} = w_{t,i} - \frac{\alpha}{\sqrt{G_{t,i} + \epsilon}} g_{t,i}$$

Adadel

$$E[\Delta_w^2]_t = \gamma E[\Delta_w^2]_{t-1} + (1 - \gamma) \Delta_w^2$$

$$\Delta w_t = \frac{\sqrt{E[\Delta_w^2]_t + \epsilon}}{\sqrt{G_{t,i} + \epsilon}} g_t$$

# Optimizers

ADAM: decaying average of the past squared gradients  
and momentum

RMSprop /

Adadelta

$$g_{t,i} = \nabla_w E(w_{t,i})$$

$$G_{t+1,i} = \gamma G_{t,i} + (1 - \gamma) g_{t,i} \odot g_{t,i}$$





$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$$

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$$

$$\hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}$$

$$w_{t+1,i} = w_{t,i} - \frac{\alpha}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t$$




Which optimizer is the best?

# **Convolutional | Neural Networks**




# Images are a series of Pixel Values



Grayscale images:  
0=Black  
255 = White

Spatial locality structure


# Handling images with Neural Networks



Works well for simple images, but fails when there are more complex patterns in the image


# Local receptive fields

Make connections in small, localized regions of the input image



# Local receptive fields

Slide the local receptive field over by one (or more) pixel and repeat



# The convolution operation

Image

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

1	0	1
0	1	0
1	0	1

Filter/  
Feature detector

1. Pointwise multiply
2. Add results
3. Translate filter

1 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	0	0
0 <sub>x0</sub>	1 <sub>x1</sub>	1 <sub>x0</sub>	1	0
0 <sub>x1</sub>	0 <sub>x0</sub>	1 <sub>x1</sub>	1	1
0	0	1	1	0
0	1	1	0	0

Image



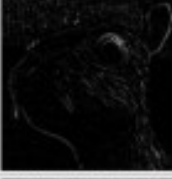


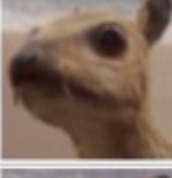

4		

Convolved  
Feature

# Filters

Original Image




Operation	Filter	Convolved Image
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	
Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Box blur (normalized)	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	
Gaussian blur (approximation)	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	




# Stride

7 x 7 Input Volume




5 x 5 Output Volume



Stride 1

7 x 7 Input Volume

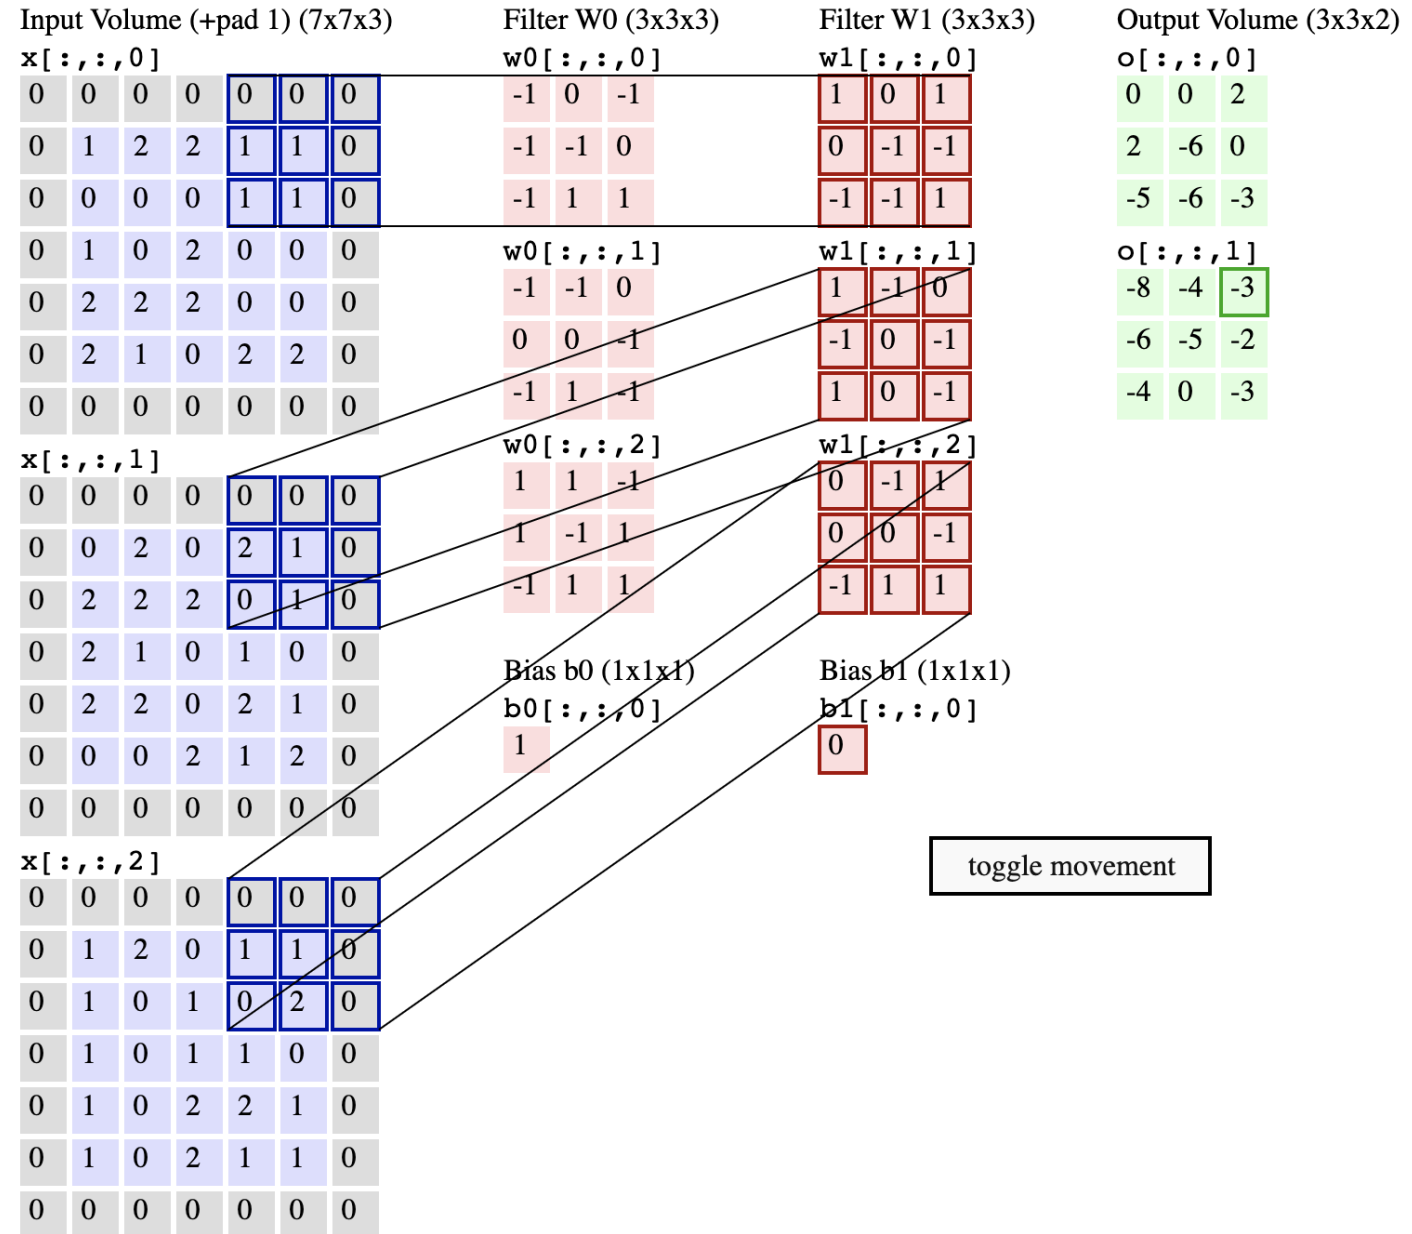


3 x 3 Output Volume



Stride 2

# CNN over the image channels




# Kernels



Example filters learned by Krizhevsky et al. Each of the 96 filters shown here is of size  $[11 \times 11 \times 3]$ , and each one is shared by the  $55 \times 55$  neurons in one depth slice. Notice that the parameter sharing assumption is relatively reasonable: If detecting a horizontal edge is important at some location in the image, it should intuitively be useful at some other location as well due to the translationally-invariant structure of images. There is therefore no need to relearn to detect a horizontal edge at every one of the  $55 \times 55$  distinct locations in the Conv layer output volume.

# Pooling




Pooling layer downsamples the volume spatially, independently in each depth slice of the input volume. Left: In this example, the input volume of size  $[224 \times 224 \times 64]$  is pooled with filter size 2, stride 2 into output volume of size  $[112 \times 112 \times 64]$ . Notice that the volume depth is preserved. Right: The most common downsampling operation is max, giving rise to max pooling, here shown with a stride of 2. That is, each max is taken over 4 numbers (little  $2 \times 2$  square).

# Pooling layers

- Intuition: the exact location of a feature isn't as important as its rough location
  - Helps prevent overfitting
- Reduces the number of parameters needed in later layers
- $L_2$  pooling is also common ( $L_2$  norm)

# Fully connected layer to combine

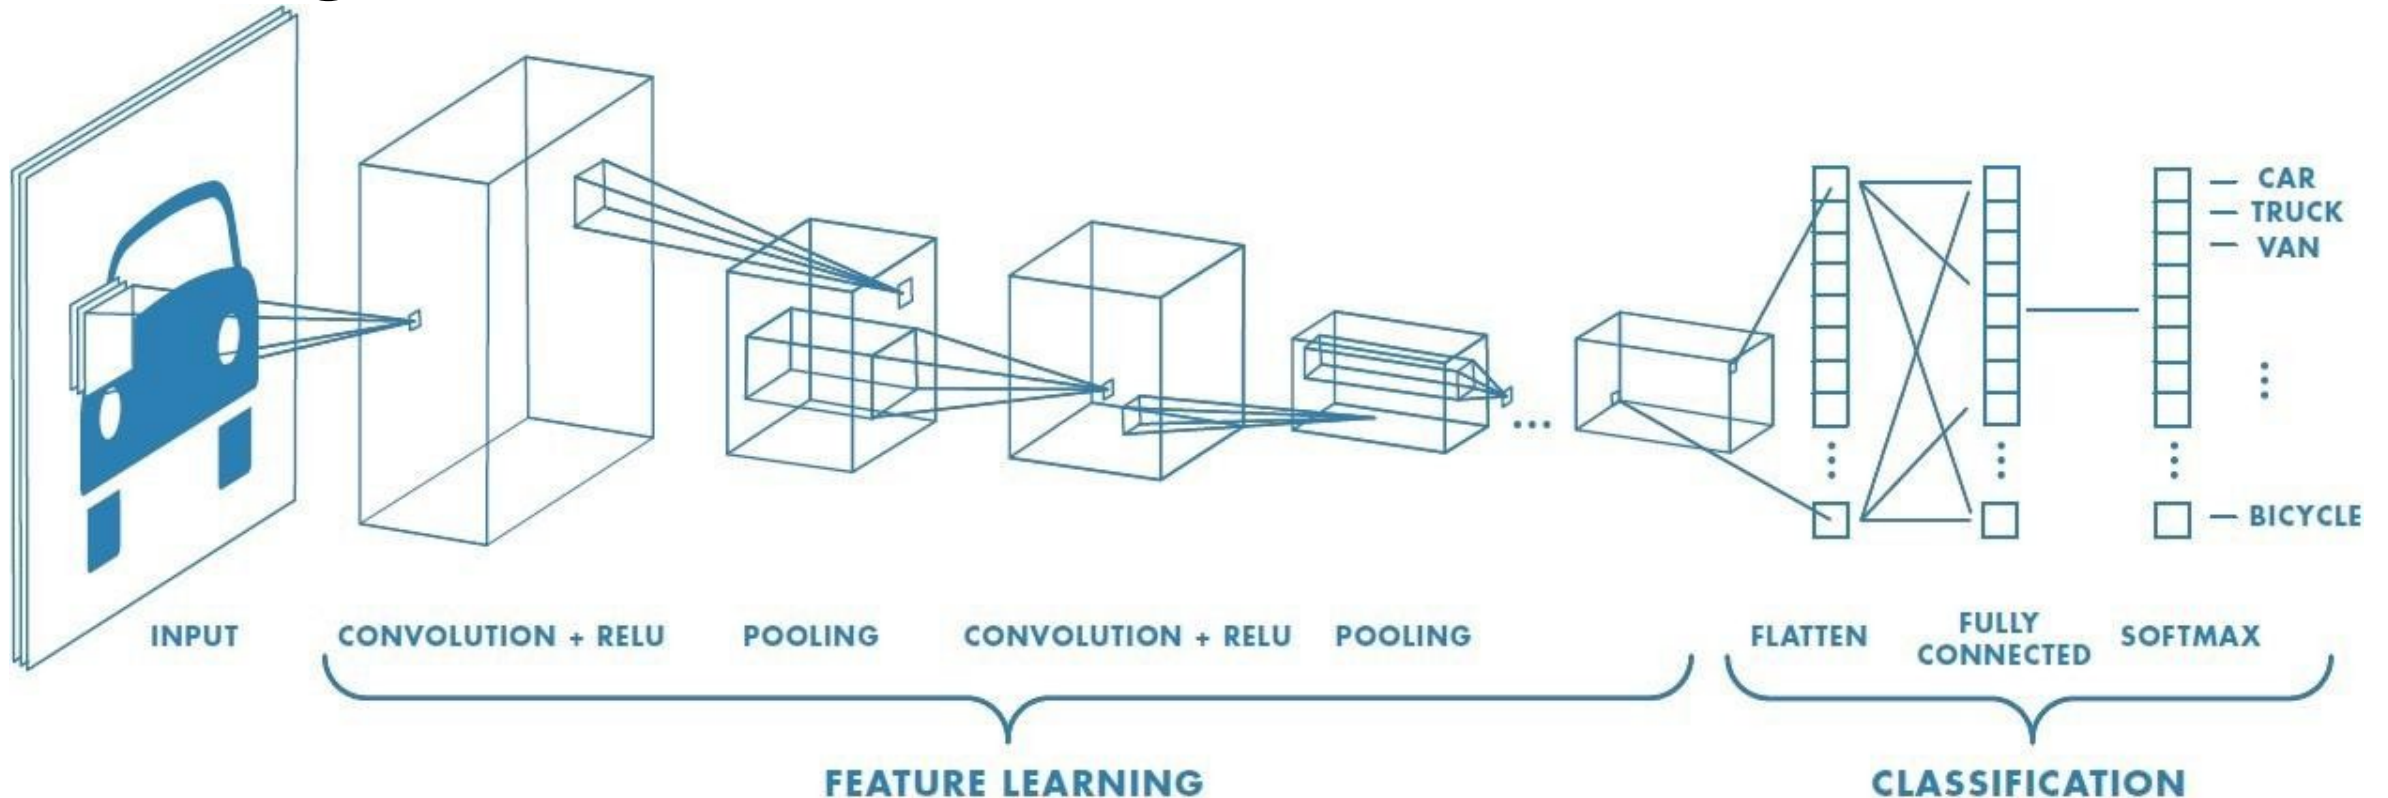
- Convolutional layers detected features
- Pooling layers reduced complexity
- Now we have a set of feature maps



sify




# Image Classification with CNN




- CONV and POOL layers output high-level features of input
- Fully connected layer uses these features for classifying input image
- Express output as probability of image belonging to a particular class

$$\text{softmax}(y_i) = \frac{e^{y_i}}{\sum_j e^{y_j}}$$

# CNN and brain architecture



Hubel and Wiesel,  
1959-1968





Fukushima,  
1980



Brain “**inspired**” model

# Which pixels matter: Saliency via Occlusion

Mask part of the image before feeding to CNN,  
check how much predicted probabilities change



$$P(\text{elephant}) = 0.95$$



$$P(\text{elephant}) = 0.75$$





Zeiler and Fergus, "Visualizing and Understanding Convolutional Networks", ECCV 2014

[Boat image](#) is [CC0 public domain](#)  
[Elephant image](#) is [CC0 public domain](#)  
[Go-Karts image](#) is [CC0 public domain](#)



# Which pixels matter: Saliency via Occlusion


Mask part of the image before feeding to CNN,  
check how much predicted probabilities change




Zeiler and Fergus, "Visualizing and Understanding Convolutional Networks", ECCV 2014

[Boat image is CC0 public domain](#)  
[Elephant image is CC0 public domain](#)  
[Go-Karts image is CC0 public domain](#)


schooner



African elephant, *Loxodonta africana*




go-kart



# Which pixels matter: Saliency via Backprop

Forward pass: Compute probabilities




Simonyan, Vedaldi, and Zisserman, "Deep Inside Convolutional Networks: Visualising Image Classification Models and Saliency Maps", ICLR Workshop 2014.


Figures copyright Karen Simonyan, Andrea Vedaldi, and Andrew Zisserman, 2014; reproduced with permission.

# Which pixels matter: Saliency via Backprop

Forward pass: Compute probabilities







Compute gradient of (unnormalized) class score with respect to image pixels, take absolute value and max over RGB channels





# Saliency Maps



Simonyan, Vedaldi, and Zisserman, "Deep Inside Convolutional Networks: Visualising Image Classification Models and Saliency Maps", ICLR Workshop 2014.  
Figures copyright Karen Simonyan, Andrea Vedaldi, and Andrew Zisserman, 2014; reproduced with permission.



Time for a quiz and tutorial!



[https://tinyurl.com/  
GeoComp2023](https://tinyurl.com/GeoComp2023)