

Using GPT as your coding partner

Antonio Fonseca - Matera 2023

Agenda

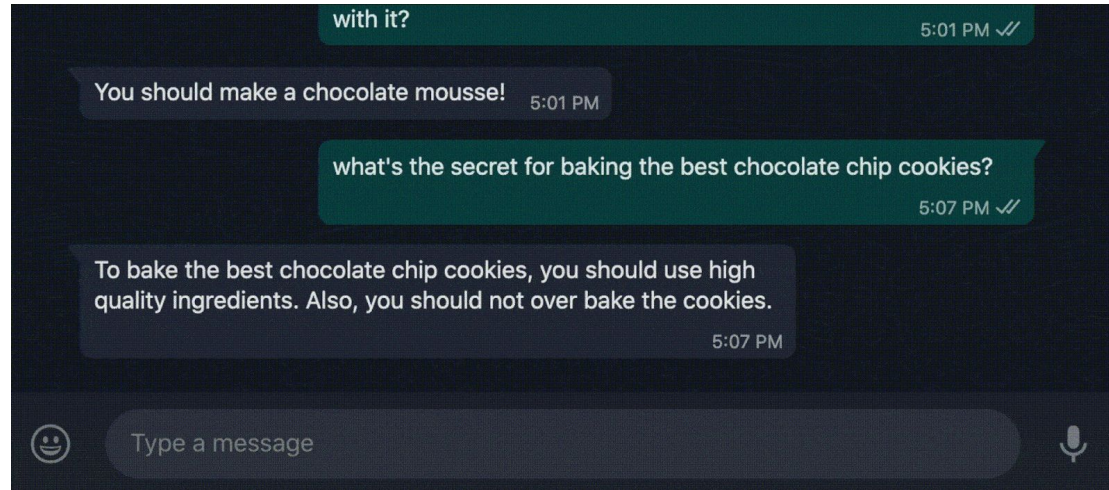
- What is GPT?
- How does it work?
- Overview:
 - Pretraining
 - Prompting
 - Inference
- Tutorial

Generative AI and large datasets

Dall-E: generate image from prompt



ChatGPT: conversational agent



Introduction to GPT

What is GPT?

- GPT stands for Generative Pre-trained Transformer.
- It is a language model developed by OpenAI that uses deep learning to generate human-like text.
- Tasks:
 - Language translation
 - Text summarization
 - PowerPoint presentations (used to create these slides)
 - Coding

How does GPT work?

The basics

- GPT uses a neural network to analyze large amounts of text data and learn patterns in language.
- Once trained, it can generate text based on a given prompt.
- The more data it is trained on, the better it becomes at generating human-like text.

GPT Overview

- Estimated cost of pretraining GPT-3: 355 GPU years (NVIDIA Tesla V100 GPU) and cost \$4.6m.
- 175B parameters
- Massive training set
- Up to 2048 tokens (words) for context
- Main steps:
 - Pretraining,
 - Prompting
 - Inference


GPT-3 training data^{[1]:9}

Dataset	# tokens	Proportion within training
Common Crawl	410 billion	60%
WebText2	19 billion	22%
Books1	12 billion	8%
Books2	55 billion	8%
Wikipedia	3 billion	3%


Step 1

Collect demonstration data, and train a supervised policy.


A prompt is sampled from our prompt dataset.



A labeler demonstrates the desired output behavior.




This data is used to fine-tune GPT-3 with supervised learning.



Step 2

Collect comparison data, and train a reward model.


A prompt and several model outputs are sampled.



A labeler ranks the outputs from best to worst.



This data is used to train our reward model.




Step 3

Optimize a policy against the reward model using reinforcement learning.

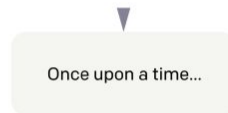
A new prompt is sampled from the dataset.





The policy generates an output.



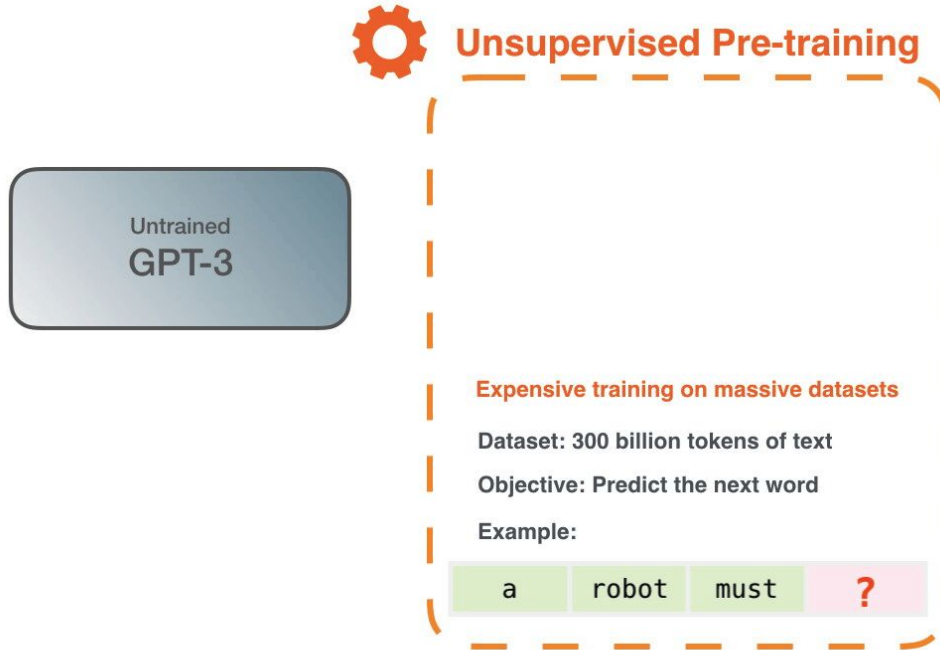
The reward model calculates a reward for the output.



The reward is used to update the policy using PPO.



GPT Pretraining



GPT Pretraining

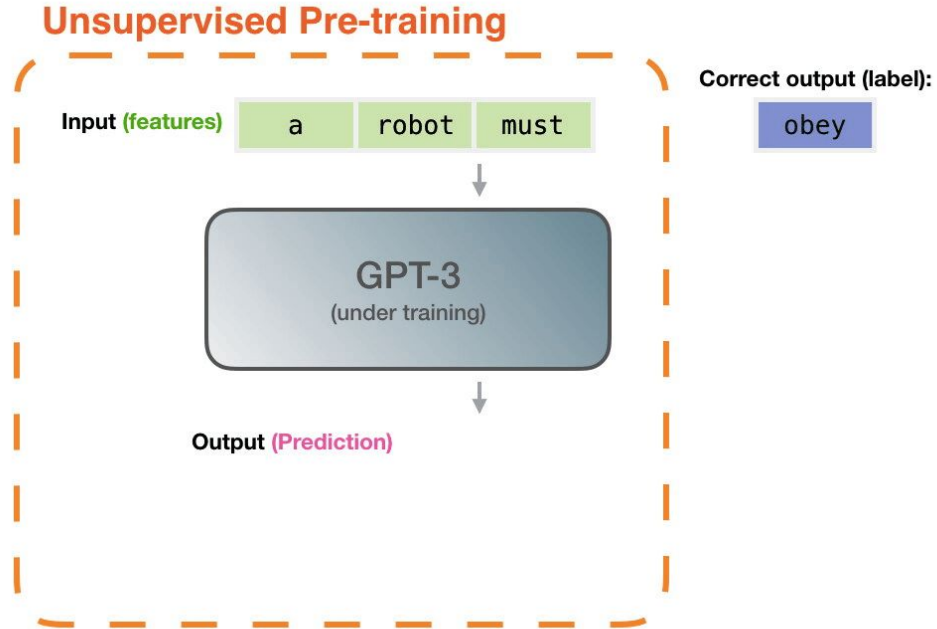
Text: Second Law of Robotics: A robot must obey the orders given it by human beings



Generated training examples

Example #	Input (features)	Correct output (labels)
1	Second law of robotics :	a
2	Second law of robotics : a	robot
3	Second law of robotics : a robot	must
...		

GPT Pretraining




Prompting

Input Prompt: Recite the first law of robotics

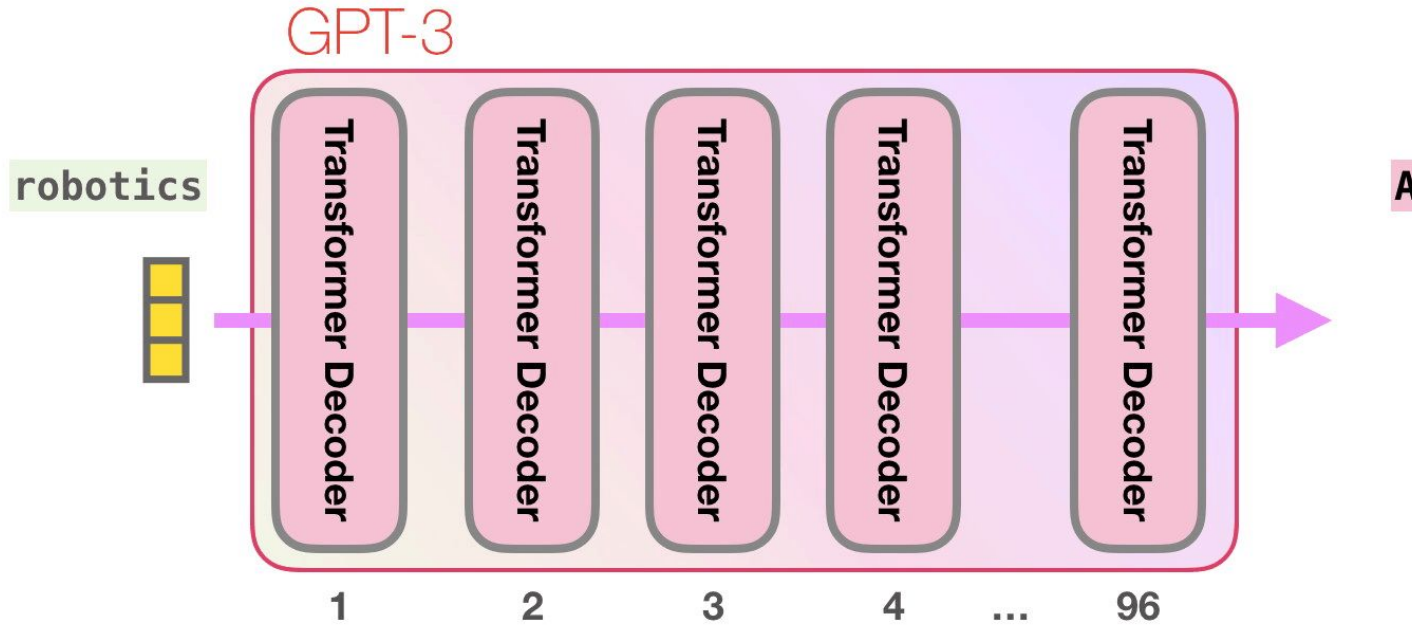


Output:

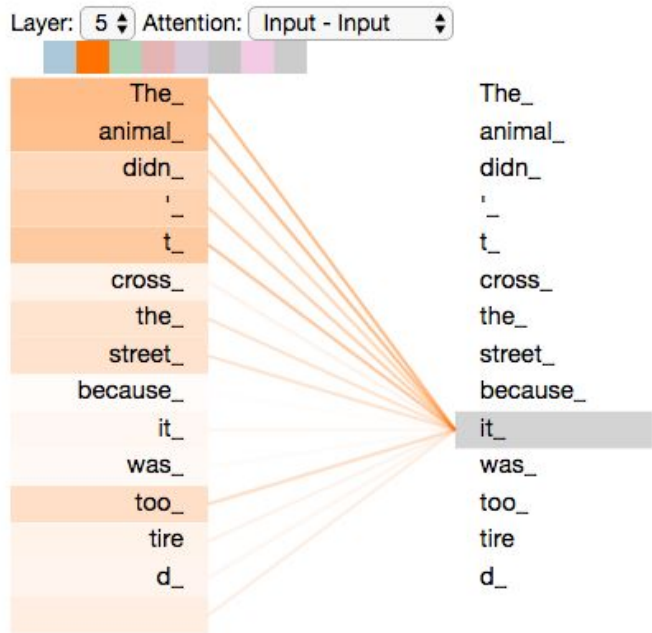
Inference: next token prediction



The “magic” (Transformers)



Transformer's secret sauce: self-attention



Example:

*"The animal didn't cross the street because **it** was too tired"*

What does "it" in this sentence refer to?

Self-attention allows it to look at other positions in the input sequence for clues that can help lead to a better encoding for this word.

GPT-4

The latest milestone

- OpenAI has recently released GPT-4, the latest version of the language model.
- GPT-4 is a large (100T parameters) multimodal model that can accept both text and images as input.
- It is capable of generating text outputs in a variety of domains, including natural language and code
- Generates code in a variety of programming languages, including Python, JavaScript, and HTML.